

# Human-Sounding AI Copy

Writing that doesn't scream ChatGPT — by Obed Cudjoe

# **Human-Sounding AI Copy**

The anti-generic playbook. Read it in an evening, use it the next morning.

# The AI Tell: Why Everyone Recognizes ChatGPT Copy

You've felt it reading an article, an ad, a LinkedIn post: a certain smoothness. Balanced sentences. A tidy structure. And absolutely nothing that sounds like a human made a choice. That's the AI tell - and once you see it, you can't unsee it.

The tell is not one thing. It's a cluster of tells. The em-dash that shows up three times per paragraph. The 'it's not just X, it's Y' construction that appears in every AI-written piece on the internet. The rhythm that never varies - long sentence, short sentence, long sentence, perfectly alternating. The total absence of anything specific: no odd detail, no unfashionable opinion, no sentence that a hundred other writers wouldn't write.

Why does this happen? Language models are trained to predict the most likely next word. The most likely word is the average word. The most likely structure is the average structure. Left to their own devices, they produce the average of all copy ever written - which is why AI copy feels like it was written by a committee of nobody.

The good news: the average is avoidable. Models are capable of surprising, specific, opinionated writing - but they need to be pushed off the average. That push comes in two forms: constraints that forbid the average moves, and inputs that give them something specific to work with. Both are in your control.

This book is the control panel. Chapter two builds the voice brief that teaches AI to sound like your brand. Chapter three is the editing pass that removes the tells from anything already written. Chapter four gives you prompt patterns for ads, emails and landing pages that start specific and stay specific.

The goal is not to disguise AI writing as human writing. The goal is writing that happens to be good - and good writing, whether a human or a model drafted it, is specific, rhythmical and a little bit strange. You're about to learn how to get there on purpose.

Action steps: (1) Find one piece of AI-generated copy you've published. (2) Circle every em-dash, every 'it's not just X' and every sentence that could appear in any article about any topic. (3) You now have your first editing list.

# The Voice Brief: Teaching AI Your Brand Voice in 10 Minutes

Most brands don't have a voice; they have a vibe. A vibe can't be pasted into a prompt. A voice can. The difference is a voice brief - a document that tells the AI exactly how your brand talks, in terms it can follow.

A voice brief has five sections, and you can write all five in ten minutes. Section one: three words that describe how you sound (e.g. 'plain, wry, direct'). Section two: three words you'd never use (e.g. 'synergy, leverage, delightful'). Section three: sentence rhythm - short or long? fragments allowed? Section four: what you're willing to say that others won't - your contrarian positions. Section five: three real samples of your best writing, pasted verbatim.

Section five is the secret weapon. A model can imitate a described voice about as well as you can imitate a described accent. But give it three examples of the actual thing - like an actor listening to a voice recording - and the imitation sharpens dramatically. Your old emails, your best tweets, your most-mailed newsletter: those are your voice samples.

Here's the ten-minute build: minutes 1-3, write the three words, the banned words and the rhythm notes. Minutes 4-6, write your two contrarian positions (if you can't think of any, your voice needs work - brands without positions sound like everyone). Minutes 7-10, find and paste the three samples. Save it as a file. That's the brief.

Now use it: at the top of every working session, paste the brief, then say 'now write [task] in this voice'. You'll still need editing - chapter three - but the output will arrive closer to your voice than anything you've gotten before.

Action steps: (1) Write your five-section voice brief in the next 10 minutes. (2) Run one ad prompt with and one without the brief. (3) Compare - and never run without again.

## The De-AI Editing Pass: 3 Rounds to Human

Editing AI copy is not fixing typos. It's a three-round pass that strips the average and restores the specific. Each round takes about five minutes per page, and each targets one layer of the tell.

Round one: structure. Read the copy aloud. Mark every sentence you would never say in conversation. Rewrite those as you'd actually speak - fragments, contractions, the occasional sentence that starts with 'and'. The ear is a brutal editor; if it sounds like a voiceover, it sounds like AI. This round fixes rhythm.

Round two: specificity. The AI wrote in categories; you edit into particulars. Replace every noun with a real thing: 'customers' becomes 'boutique owners in Osu'; 'results' becomes 'bookings up 30% in six weeks'. Add one detail only you know - a client story, a number from your own dashboard, an odd fact. Specificity is the fastest way to sound human, because AI can't invent your history and won't.

Round three: opinion. Find the places where the copy hedges - 'arguably', 'can help', 'may be'. Humans don't hedge when they believe something. Replace hedges with positions: 'this works' instead of 'this can work'. If the copy has no positions at all, add one: the thing you believe that your competitors won't say. Opinion is the final tell-remover.

Three rounds, fifteen minutes a page, and the copy reads like a human who cares. Over time you'll internalize the rounds and do them in one pass - but always do the pass. The moment you skip editing is the moment your readers feel the tell.

Action steps: (1) Take the copy you circled in chapter one. (2) Run the three rounds on one page. (3) Before-and-after: read both aloud. Keep the after.

## Prompt Patterns for Ads, Emails & Landing Pages

This chapter is the pattern library: four prompt structures that consistently produce copy closer to human, plus the constraints that make them work. Use them with your voice brief from chapter two.

Pattern one: the outcome-first ad. Prompt structure: audience, outcome, proof, then 'write 5 headlines under 8 words, each promising one outcome'. The constraint 'one outcome per headline' forces specificity; the proof line gives the model real material instead of averages. This pattern produces ads that read like a confident person made a claim.

Pattern two: the objection-crushing email. Prompt structure: the offer, the list's last objection (write it as a real quote), then 'write a 200-word email that opens with the objection, answers it with one proof point, and closes with one CTA'. Opening with the objection makes the email feel like a conversation, not a broadcast - the single most human move in email marketing.

Pattern three: the story landing page. Prompt structure: the customer, the moment they nearly gave up, the solution, the result, then 'write a landing page hero, subhead and 3 sections in the shape of a story: before, turning point, after'. Stories are specific by nature - they contain people, times and places - which is exactly what AI copy lacks by default.

Pattern four: the banned-phrase list. Every prompt in this book works better with a short list of what the AI may not say. Start with: 'I hope this finds you well', 'in today's fast-paced world', 'game-changer', 'unlock the power', 'seamless', 'elevate'. Six bans remove half the tells instantly.

The common thread: every pattern gives the AI something real to work with (proof, objections, stories) and something to avoid (the banned list). Real material plus constraints equals copy that doesn't sound like the average of the internet - because it isn't.

Action steps: (1) Rewrite one live ad using pattern one. (2) Rewrite one email using pattern two. (3) Run both through the three editing rounds. Ship the better versions.

## **Your Voice Library: Checklists & 30-Day Plan**

The final chapter is about compounding. Every piece of copy you write - and every prompt that produced it - is an asset. The brands that sound most human aren't the most talented; they're the most consistent. Consistency comes from a system, and your system is the voice library.

The library has four shelves. Shelf one: your voice brief (chapter two), updated monthly. Shelf two: your wins - every piece of copy that performed, with the prompt that produced it. Shelf three: your banned lists - the phrases your brand will never say, growing as you spot new tells. Shelf four: your proof bank - results, testimonials and numbers, ready to paste into any prompt.

Here's the 30-day plan. Week one: build the brief, run every prompt with it, start the wins shelf. Week two: add the three-round editing pass to every draft; move your best two pieces to the wins shelf. Week three: build the proof bank - dig out five real results and paste them everywhere. Week four: review the month - which prompts produced the most keepers? Promote those to your favourites and delete the rest.

By day 30 you won't recognize your AI output. The same model, the same tools - but now fed with your voice, your proof and your constraints, and edited with your ear. That's the whole game: not better AI, a better system around it.

One caution as you go: never publish AI copy unedited, no matter how good the prompt. The tells hide in plain sight, and your readers have been trained to see them. The system exists to make the editing fast - not to make it unnecessary.

Action steps: (1) Create the four shelves today. (2) Add your first win tonight. (3) Revisit this page in 30 days - and count how many 'wins' you've accumulated.